

YOLO Detection PYTORCH Model Inference Report

Backend Engine: PyTorch

Phase 1: Configuration & System Metadata

Model Name	yolo26n.pt
Input Dimensions	640x640
Source Images Directory	/workspace/Images
Processed Images Output	/workspace/PT_Output
OS	Ubuntu 24.04.4 LTS
Python Version	3.12.3
PyTorch Version	2.13.0+cu130
Ultralytics Version	8.4.75
Inference Device	NVIDIA RTX A4000
CUDA Version	13.0
Total Parameters	2,408,932
File Size	5.29 MB
Number of Classes	80
Precision Mode	FP32

Model Trained Class Names:

person, bicycle, car, motorcycle, airplane, bus, train, truck, boat, traffic light, fire hydrant, stop sign, parking meter, bench, bird, cat, dog, horse, sheep, cow, elephant, bear, zebra, giraffe, backpack, umbrella, handbag, tie, suitcase, frisbee, skis, snowboard, sports ball, kite, baseball bat, baseball glove, skateboard, surfboard, tennis racket, bottle, wine glass, cup, fork, knife, spoon, bowl, banana, apple, sandwich, orange, broccoli, carrot, hot dog, pizza, donut, cake, chair, couch, potted plant, bed, dining table, toilet, tv, laptop, mouse, remote, keyboard, cell phone, microwave, oven, toaster, sink, refrigerator, book, clock, vase, scissors, teddy bear, hair drier, toothbrush

Phase 2: Inference Performance Metrics

Total Images Processed	5
Total Run Time	294.34 ms
Preprocess Latency	1.88 ms
Inference Latency	55.29 ms
Postprocess Latency	1.70 ms
Peak Memory Usage	90.00 MB
Throughput	16.99 Images/sec

Phase 3: Detailed Predictions Summary

Image File	Detected Objects (Confidence)	Bounding Boxes (xyxy)
Cat_November_2010-1a.jpg.webp	cat (91.9%)	[208, 146, 888, 1221]
images.jpeg	dog (94.3%)	[86, 148, 422, 514]
images1.jpeg	zebra (91.9%)	[27, 171, 587, 409]
istockphoto-1189925691-612x612.jpg	bus (94.8%)	[73, 81, 577, 350]
tigermain.jpg	zebra (95.9%)	[680, 420, 1873, 1090]