

YOLO Detection PYTORCH Model Inference Report

Backend Engine: PyTorch

Phase 1: Configuration & System Metadata

Model Path	/workspace/Detection_Models/yolo26n.pt
Input Dimensions	640x640
Source Images Directory	/workspace/Images
Processed Images Output	/workspace/PT_Output
OS	Linux 6.17.0-40-generic
Python Version	3.12.3
PyTorch Version	2.13.0+cu130
Ultralytics Version	8.4.75
Inference Device	NVIDIA RTX A4000
CUDA Version	13.0
Model Architecture	YOLO (Detect)
Total Parameters	2,408,932
File Size	5.29 MB
Number of Classes	80
Precision Mode	FP32

Model Trained Class Names:

person, bicycle, car, motorcycle, airplane, bus, train, truck, boat, traffic light, fire hydrant, stop sign, parking meter, bench, bird, cat, dog, horse, sheep, cow, elephant, bear, zebra, giraffe, backpack, umbrella, handbag, tie, suitcase, frisbee, skis, snowboard, sports ball, kite, baseball bat, baseball glove, skateboard, surfboard, tennis racket, bottle, wine glass, cup, fork, knife, spoon, bowl, banana, apple, sandwich, orange, broccoli, carrot, hot dog, pizza, donut, cake, chair, couch, potted plant, bed, dining table, toilet, tv, laptop, mouse, remote, keyboard, cell phone, microwave, oven, toaster, sink, refrigerator, book, clock, vase, scissors, teddy bear, hair drier, toothbrush

Phase 2: Inference Performance Metrics

Total Images Processed	5
Total Run Time	906.16 ms
Preprocess Latency	1.81 ms per image
Inference Latency	52.84 ms per image
Postprocess Latency	1.65 ms per image
Throughput	5.52 Images/sec
Peak Memory Usage	69.19 MB

Phase 3: Detailed Predictions Summary

Image File	Detected Objects (Confidence)	Bounding Boxes (xyxy)
Cat_November_2010-1a.jpg.webp	cat (91.9%)	[208, 146, 888, 1221]
images.jpeg	dog (94.3%)	[86, 148, 422, 514]
images1.jpeg	zebra (91.9%)	[27, 171, 587, 409]
istockphoto-1189925691-612x612.jpg	bus (94.8%) bus (36.5%)	[73, 81, 577, 350] [573, 196, 611, 288]
tigermain.jpg	zebra (95.9%)	[680, 420, 1873, 1090]